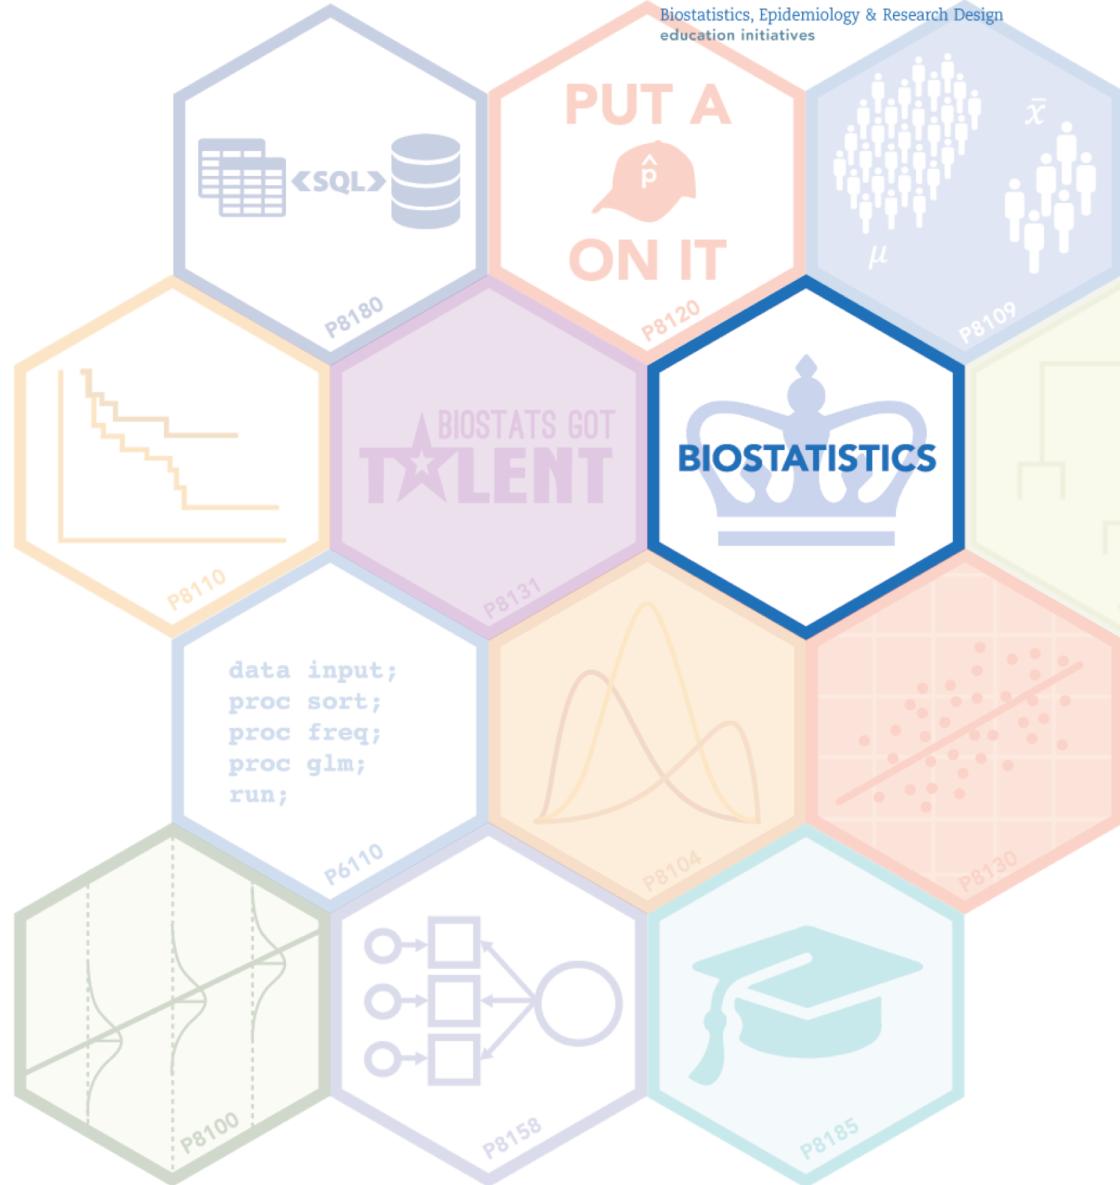


GETTING STARTED WITH R (PART 2)

Christine Mauro, PhD

November 26, 2018



Last Time

- Overview of R and R Studio
- Importing Data
 - Read CSV files using `readr` package
- Examining Data Attributes
 - Data structure, type and dimensionality
- Manipulating Data (Data Wrangling)
 - Select, Filter, Mutate, Arrange
 - Stacking and Merging

Today's Outline

- Descriptive Statistics
 - Continuous/Categorical data
- Data Visualization
 - Histogram
 - Box-plot
 - Scatterplot
- Basic Hypothesis Testing
 - T-tests and ANOVA
 - Chi-squared and Fisher's Exact test

Application

- Risk Factors Associated with Low Birthweight: `lowbwt_ALL.csv`
- The data on 189 births were collected at Baystate Medical Center, Springfield, Mass. during 1986. The dataset contains an indicator of low infant birth weight as a response and several risk factors associated with low birth weight. The actual birth weight is also included in the dataset.
- The dataset consists of the following 10 variables:
 - `low`: indicator of birth weight less than 2.5kg
 - `age`: mother's age in years
 - `lwt`: mother's weight in pounds at last menstrual period
 - `race`: mothers race ("white", "black", "other")
 - `smoke`: smoking status during pregnancy (yes/no)
 - `ht`: history of hypertension (yes/no)
 - `ui`: presence of uterine irritability (yes/no)
 - `ftv`: physician visit during the first trimester (yes/no)
 - `ptl`: previous premature labor (yes/no)
 - `bwt`: birth weight in grams

Let's Get Started

- Step 1: Open your R project from last time (double click file to open)
- Step 2: Create a new script and save.
- Step 3: Load in lowbwt_ALL.csv and packages.

```
1 - #####  
2 # November 26, 2018  
3 # Christine Mauro  
4 #  
5 # Getting Started with R - Part 2  
6 - #####  
7  
8 library(dplyr)  
9 library(readr)  
10  
11 - ##### Loading Data #####  
12  
13  
14 lowbirth = read_csv(file = "./lowbwt_All.csv")  
15 names(lowbirth)  
16 lowbirth
```

DESCRIPTIVE STATISTICS

Descriptive Stats

- Describe the basic features of the data
- Continuous Variables
 - Measures of central tendency (e.g., mean, median)
 - Measures of variability/spread (e.g., standard deviation, interquartile range, range)
- Categorical Variables
 - Counts and Percentages

skimr package

library(skimr)

```
> skim(lowbirth)
Skim summary statistics
n obs: 189
n variables: 10

— Variable type:character —————
variable missing complete n min max empty n_unique
  race      0       189 189   5   5      0       3

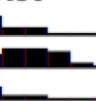
— Variable type:integer —————
variable missing complete n     mean      sd    p0    p25    p50    p75    p100 hist
  age      0       189 189 23.24  5.3   14    19    23    26    45
  bwt      0       189 189 2944.66 729.02 709 2414 2977 3475 4990
  ftv      0       189 189   0.47   0.5    0    0    0    1    1
  ht       0       189 189   0.063  0.24    0    0    0    0    1
  low      0       189 189   0.31   0.46    0    0    0    1    1
  lwt      0       189 189 129.69 30.65   80   110   121   140   250
  ptl      0       189 189   0.16   0.37    0    0    0    0    1
  smoke    0       189 189   0.39   0.49    0    0    0    1    1
  ui       0       189 189   0.15   0.36    0    0    0    0    1
```

skimr updated

```
> lowbirth2 <- mutate_at(lowbirth, vars(race, ftv, ht, low, ptl, smoke, ui), as.factor)
> skim(lowbirth2)
Skim summary statistics
n obs: 189
n variables: 10

— Variable type:factor —
variable missing complete   n n_unique          top_counts ordered
  ftv      0     189 189      2             0: 100, 1: 89, NA: 0 FALSE
    ht      0     189 189      2             0: 177, 1: 12, NA: 0 FALSE
    low     0     189 189      2             0: 130, 1: 59, NA: 0 FALSE
    ptl     0     189 189      2             0: 159, 1: 30, NA: 0 FALSE
    race    0     189 189      3 whi: 96, oth: 67, bla: 26, NA: 0 FALSE
    smoke   0     189 189      2             0: 115, 1: 74, NA: 0 FALSE
    ui      0     189 189      2             0: 161, 1: 28, NA: 0 FALSE

— Variable type:integer —
variable missing complete   n      mean       sd    p0    p25    p50    p75   p100 hist
  age      0     189 189  23.24    5.3   14    19    23    26    45
  bwt      0     189 189 2944.66 729.02 709 2414 2977 3475 4990
  lwt      0     189 189 129.69 30.65  80   110   121   140   250
```



> |

Summary Function (base R)

```
> summary(lowbirth2)
   low          age         lwt          race        smoke       ht         ui        ftv        ptl        bwt
0:130  Min.    :14.00  Min.    :80.0  black:26  0:115  0:177  0:161  0:100  0:159  Min.    :709
1: 59  1st Qu.:19.00  1st Qu.:110.0 other:67  1: 74  1: 12  1: 28  1: 89  1: 30  1st Qu.:2414
                  Median :23.00  Median :121.0 white:96
                  Mean   :23.24  Mean   :129.7
                  3rd Qu.:26.00  3rd Qu.:140.0
                  Max.   :45.00  Max.   :250.0
```

Descriptive Stats: Continuous Variables

- Base R functions:

mean(mydata)	Mean of all numeric variables
mean(mydata\$myvar)	Mean of a selected numeric variable from the dataset
median(mydata\$myvar)	Median: the 50 th percentile
var(mydata\$myvar)	Variance
sd(mydata\$myvar)	Standard Deviation
min(mydata\$myvar)	Minimum value
max(mydata\$myvar)	Maximum value
range(mydata\$myvar)	Range: Min-Max
quantile(mydata\$myvar)	Quartiles; Interquartile Range: 25 th – 75 th percentiles

```
> mean(lowbirth2$bwt)
[1] 2944.656
> sd(lowbirth$bwt)
[1] 729.0224
> quantile(lowbirth2$bwt, c(.25, .75))
 25% 75%
2414 3475
>
```

Summarize function (dplyr)

- Similar to summary function, but stores results as a tibble (data set).
 - Useful for later calculations
 - Can use with “group_by” function

```
> summarize(lowbirth2, mean_bwt = mean(bwt), median_bwt = median(bwt), sd_bwt = sd(bwt))
# A tibble: 1 x 3
  mean_bwt median_bwt sd_bwt
    <dbl>      <int>   <dbl>
1    2945.       2977    729.

> as.data.frame(summarize(lowbirth2, mean_bwt = mean(bwt), median_bwt = median(bwt), sd_bwt = sd(bwt)))
  mean_bwt median_bwt   sd_bwt
1 2944.656      2977  729.0224
```

Descriptive Stats: Continuous Variables

- Summary statistics for each level of another categorical variable --> use **group_by** function (**dplyr**)

Example: summary stats of birthweight ‘bwt’ by ‘race’

```
34 group_by(lowbirth2, race) %>%  
35 summarize(mean_bwt = mean(bwt), median_bwt = median(bwt), sd_bwt = sd(bwt))  
36  
  
> group_by(lowbirth2, race) %>%  
+ summarize(mean_bwt = mean(bwt), median_bwt = median(bwt), sd_bwt = sd(bwt))  
# A tibble: 3 x 4  
  race  mean_bwt median_bwt sd_bwt  
  <fct>    <dbl>     <dbl>   <dbl>  
1 black     2720.      2849     639.  
2 other     2804.      2835     721.  
3 white     3104.      3076     728.  
>
```

Descriptive Stats: Categorical Variables

- Row, column, and total frequencies
- Two- and three-way tabulations
- R functions:

```
tbl <- table(mydata$var1, mydata$var2)  
prop.table(tbl, 1)  
prop.table(tbl, 2)  
prop.table(tbl)  
xtabs(~var1+var2+var3, data=mydata)
```

Two-way table
Row proportions
Column proportions
Total proportions
3-way cross-tabulation

Let's try it

```
> tbl <- table(lowbirth2$race, lowbirth2$smoke)
> tbl

      0   1
black 16 10
other 55 12
white 44 52
> prop.table(tbl, 1)

          0         1
black 0.6153846 0.3846154
other 0.8208955 0.1791045
white 0.4583333 0.5416667
> prop.table(tbl, 2)

          0         1
black 0.1391304 0.1351351
other 0.4782609 0.1621622
white 0.3826087 0.7027027
>
```

Descriptive Stats: Categorical Variables

- 3-way tabulation

R function: `xtabs(~var1+var2+var3, data=mydata)`

Example: two-way tables of ‘race’ x ‘smoke’ stratified by the levels of history of hypertension ‘ht’

```
> xtabs(~race+smoke+ht, data=lowbirth2)
, , ht = 0

      smoke
race   0  1
  black 14  9
  other 51 12
  white 43 48

, , ht = 1

      smoke
race   0  1
  black  2  1
  other  4  0
  white  1  4
```

DATA VISUALIZATION

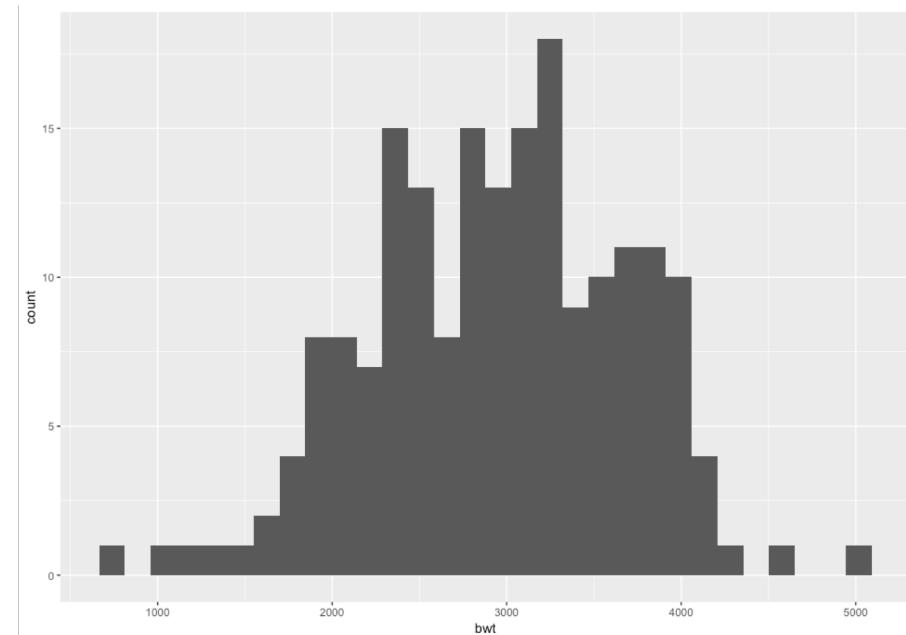
Data Visualization

- Use **ggplot2** package!
- Histogram
 - Shows the underlying frequency distribution of continuous data
- Box-plot
 - Shows the underlying distribution of continuous data based on the five number summary: min, 1st quartile, median, 3rd quartile, max.
- Scatter plot
 - Shows the relationships between two continuous (numeric) variables, each plotted of one of the axes
- Barplot
 - Useful for summarizing categorical data

Data Visualization: Histograms (ggplot2)

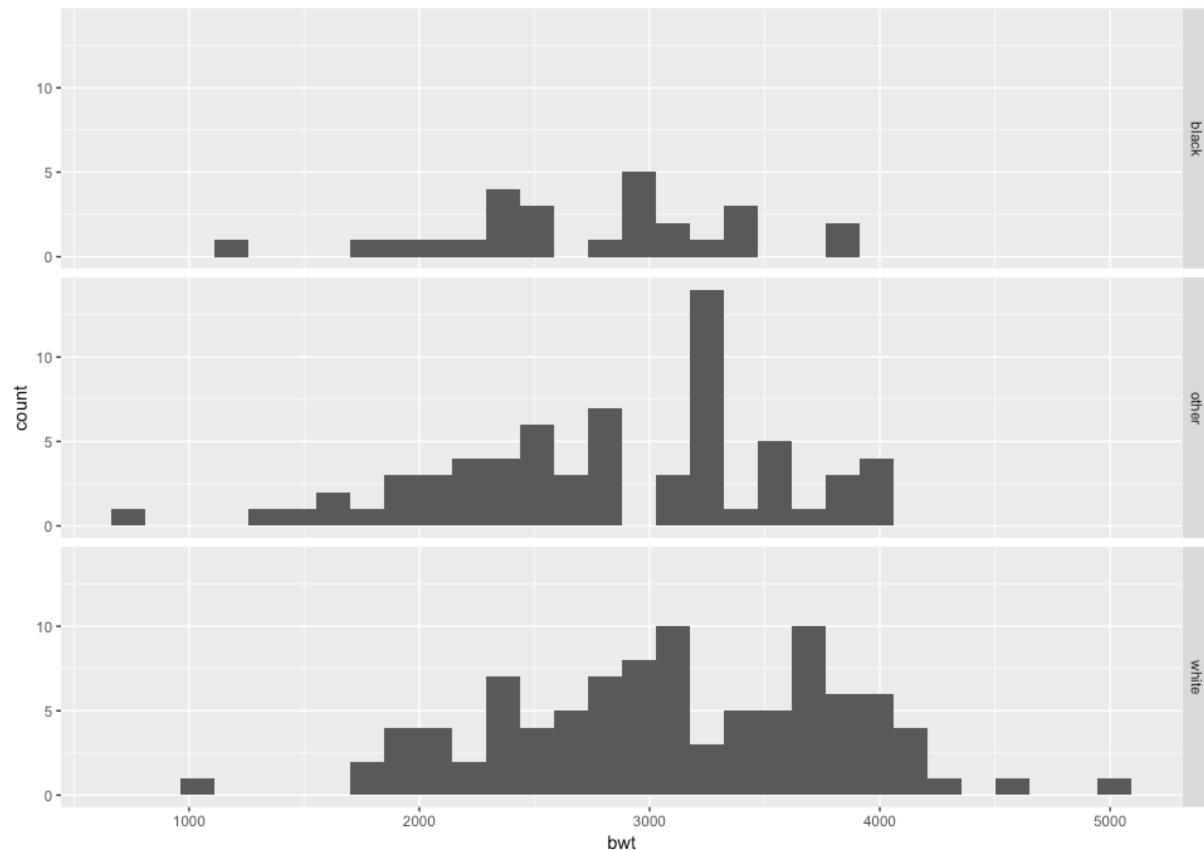
- Histogram plot of one continuous variable.

```
59 ##### Data Visualization
60
61 library(ggplot2)
62
63 #histrogram of birthweight
64 ggplot(lowbirth2, aes(x = bwt)) +
65   geom_histogram()
```



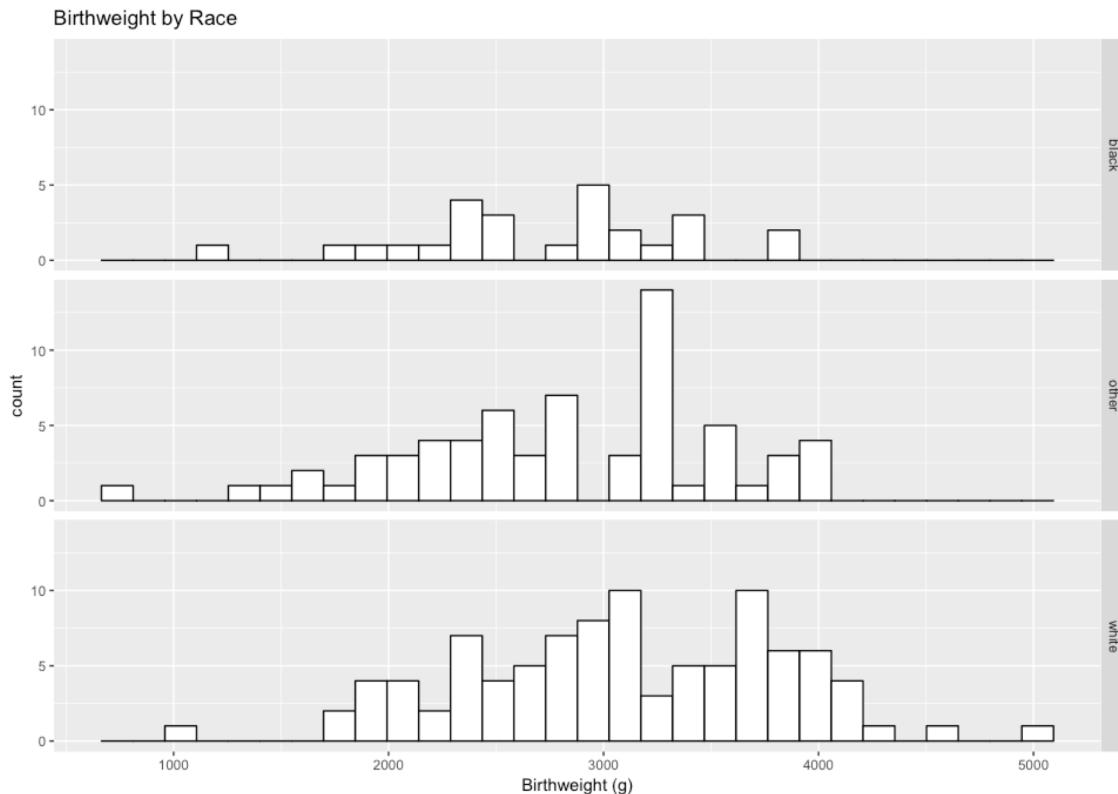
Histograms by Groups

```
68 #histogram of birthweight by race
69 ggplot(lbwirth2, aes(x = bwt)) +
70   geom_histogram() +
71   facet_grid(race ~ .)
```



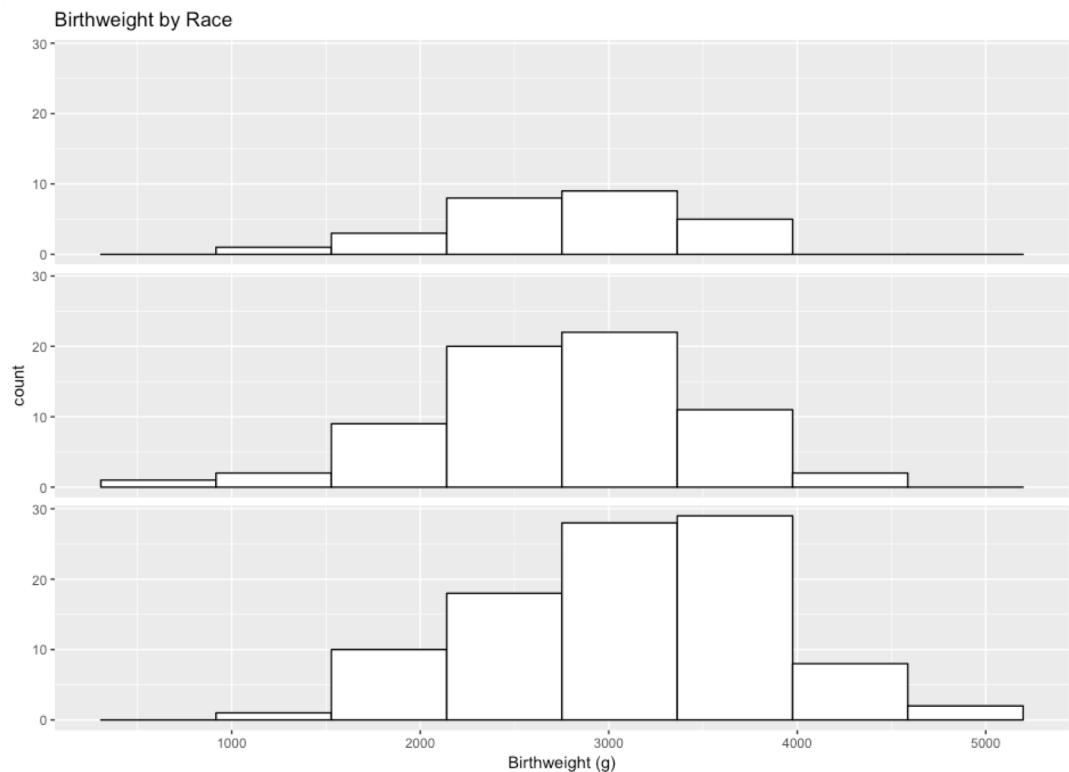
Histograms - Formatting

```
## add title and axis labels; nicer formats
ggplot(lbwirth2, aes(x = bwt)) +
  geom_histogram(colour="black", fill="white") +
  facet_grid(race ~ .) +
  ggtitle("Birthweight by Race") +
  labs(x = "Birthweight (g)")
```



Histograms – Fix Bins!

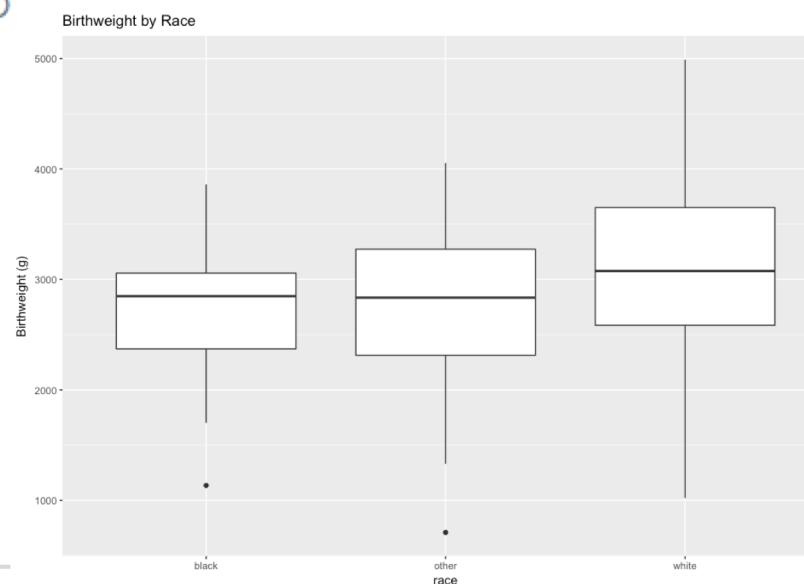
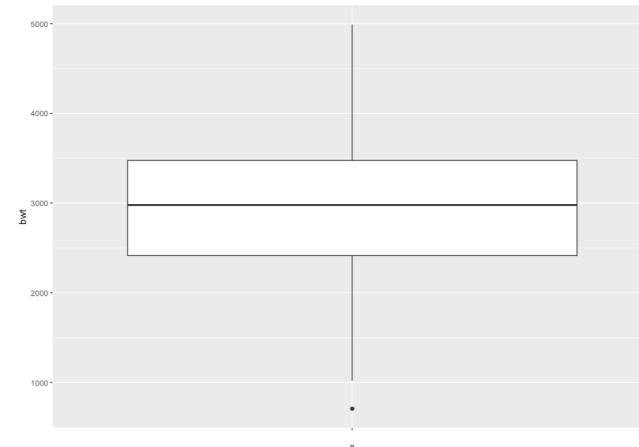
```
## add title and axis labels; nicer formats - FIX BINS
ggplot(lowbirth2, aes(x = bwt)) +
  geom_histogram(colour="black", fill="white", bins = 8) +
  facet_grid(race ~ .) +
  ggtitle("Birthweight by Race") +
  labs(x = "Birthweight (g)")
```



Data Visualization: Boxplots

Boxplot of one continuous variable; or continuous variable by categorical variable.

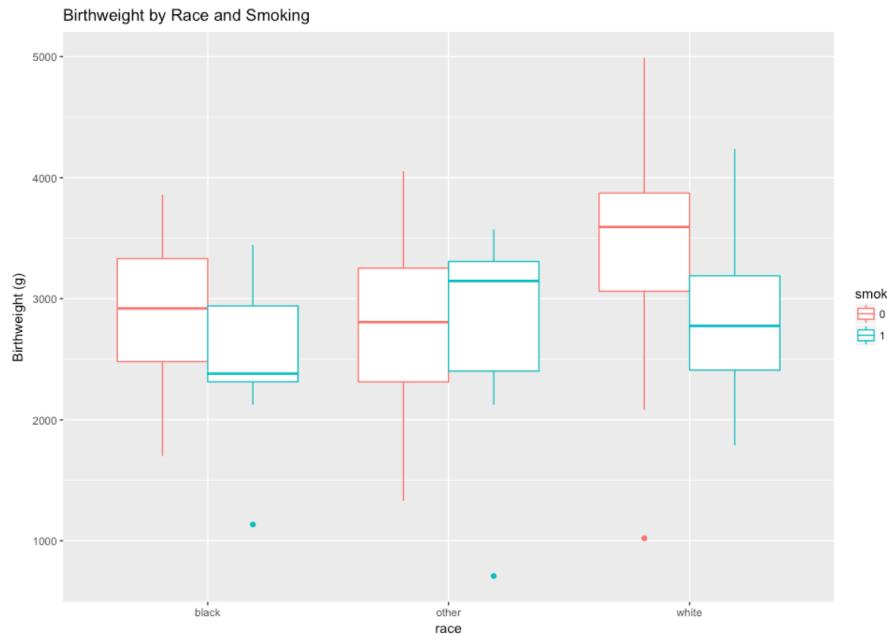
```
81 #boxplot of birthweight
82 ggplot(lowbirth2, aes(x= "", y = bwt)) +
83   geom_boxplot()
84
85
86 #boxplot of birthweight by race
87 ggplot(lowbirth2, aes(x= race, y = bwt)) +
88   geom_boxplot() +
89   ggtitle("Birthweight by Race") +
90   labs(y = "Birthweight (g)")
91
```



Data Visualization: Grouped Boxplots

A grouped boxplot is a boxplot where each category is subdivided in several groups.

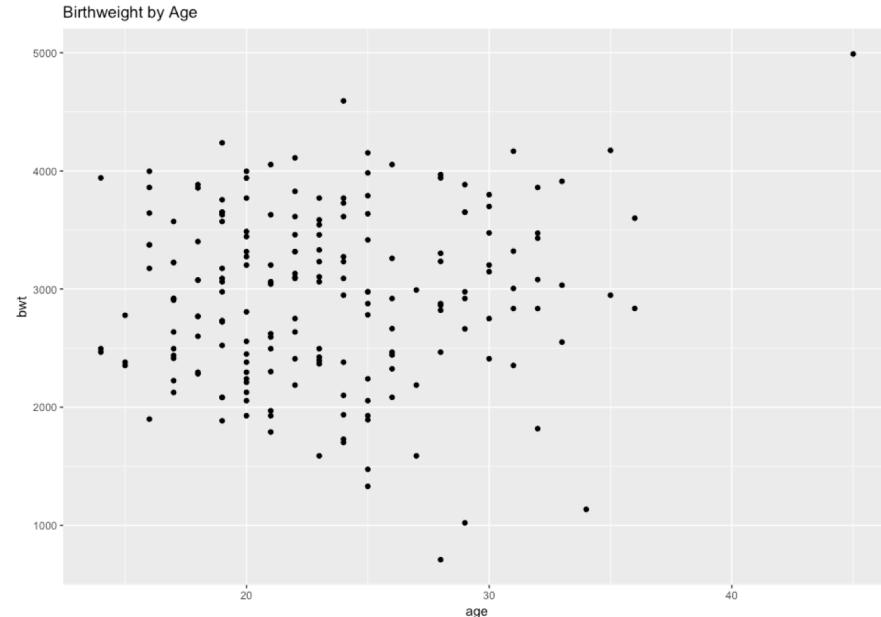
```
##  
100 #boxplot of birthweight by race and smoking status  
101 ggplot(lbwirth2, aes(x= race, y = bwt, color=smoke)) +  
102   geom_boxplot() +  
103   ggtitle("Birthweight by Race and Smoking") +  
104   labs(y = "Birthweight (g)")  
105
```



Data Visualization: Scatterplot

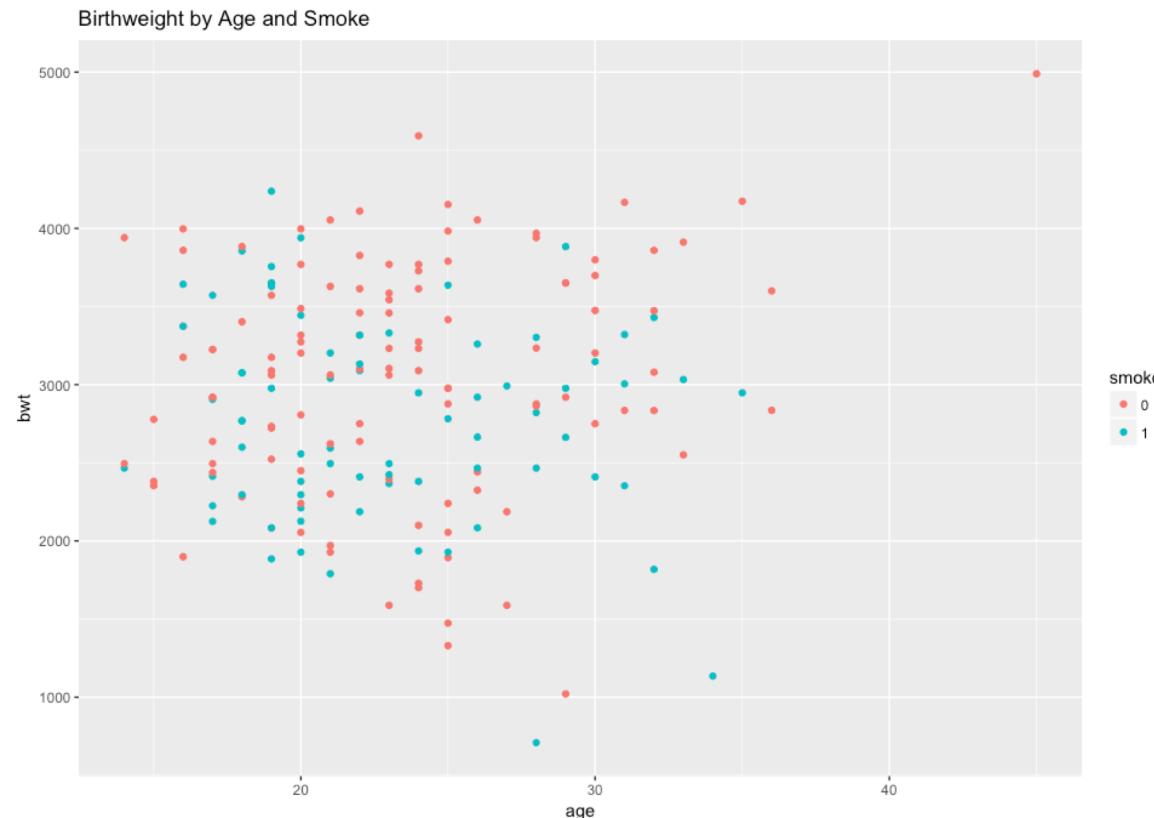
- Shows the relationship / trend between two continuous variables

```
100
101 #scatterplot of birthweight by age
102 ggplot(lbwirth2, aes(x= age, y = bwt)) +
103   geom_point() +
104   ggtitle("Birthweight by Age")
105
106
107
108
109
110
111
```



Data Visualization: Multiple Scatterplots

```
113  
114 #scatterplot of birthweight by age and smoking status |  
115 ggplot(lowbirth2, aes(x= age, y = bwt, color=smoke)) +  
116   geom_point() +  
117   ggtitle("Birthweight by Age and Smoke")  
118
```



BASIC HYPOTHESIS TESTING

Hypothesis Testing

- Hypothesis testing provides a framework for making decisions about the population based on data from a sample.

The null hypothesis (H_0)

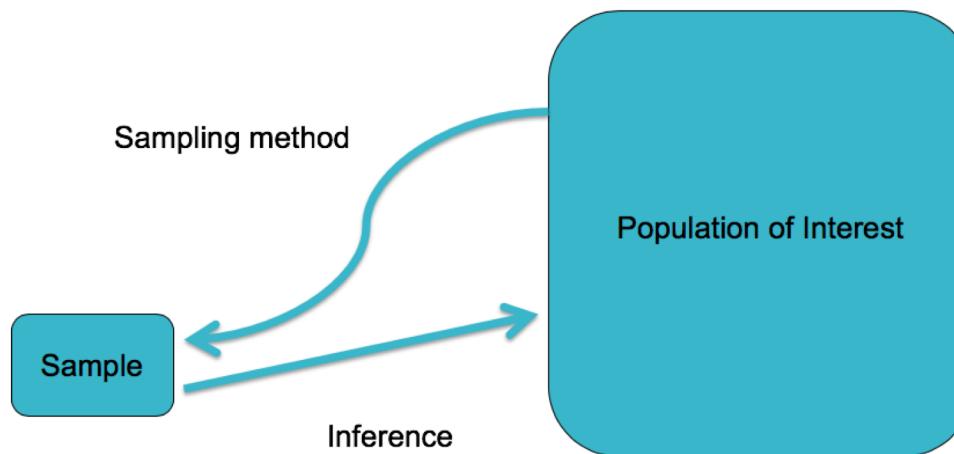
vs.

The alternative hypothesis (H_1)

- Null hypothesis is often of “no difference” in population
- Note that our decisions will always be with respect to the null hypothesis:
 - Reject the null, Fail to reject the null!

Hypothesis Testing

- Can the differences in my sample be explained by chance (i.e. sampling variability)?
 - Reject the null -> observed differences are likely not due to chance!
 - We infer from our sample back to the population:



Student's T-test

- A T-test is an analysis of one- or two-population means
 - Used for continuous (numeric) data (outcomes)
 - One-sample T-test (compare one population mean)
 - Two-sample T-test (compare two populations means)
 - Independent or paired test
- Always remember to check model assumptions before inferences
 - Normality (for small samples)
 - Histograms or QQplots/Normality tests (not covered here)
 - Independent observations within the group(s) (not repeated)

Two-Sample T-test

- Hypothesis to be tested (two-sided):

$$H_0: \mu_1 = \mu_2 \text{ vs } H_1: \mu_1 \neq \mu_2$$

- In the two-sample case, you **FIRST** need to test for the equality of variances
 - Testing the equality of variances implies testing the hypotheses:

$$H_0: \sigma_1^2 = \sigma_2^2 \text{ vs } H_1: \sigma_1^2 \neq \sigma_2^2$$

Two-Sample T-test

Example: is the birthweight of babies born to smokers significantly different than the birthweight to babies born to non-smokers?

R function to test equality of variances:

```
var.test(cont ~ binary, data=mydata)
```

R function for two-sample independent t-test:

```
t.test(cont ~ binary, data=mydata, var.equal=FALSE, paired=FALSE)
```

Options:

- Default var.equal = FALSE
 - Can be changed to TRUE if variances are unequal
- Default paired = FALSE
 - Can be changed to TRUE if data is paired, e.g., pre/post tests from same subject

Two-Sample T-test

Outcome: Continuous
Predictor: Binary

Example: is the birthweight of babies born to smokers significantly different than the birthweight to babies born to non-smokers?

```
> var.test(bwt ~ smoke, data = lowbirth2)
```

F test to compare two variances

```
data: bwt by smoke
F = 1.2993, num df = 114, denom df = 73, p-value = 0.229
alternative hypothesis: true ratio of variances is not equal to 1
95 percent confidence interval:
 0.8469514 1.9550579
sample estimates:
ratio of variances
 1.299335

> t.test(bwt ~ smoke, data=lowbirth2, var.equal= TRUE)
```

Two Sample t-test

```
data: bwt by smoke
t = 2.6336, df = 187, p-value = 0.009156
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 70.69274 492.73382
sample estimates:
mean in group 0 mean in group 1
 3054.957      2773.243
```

Not enough evidence to declare inequality of variances.

Interpretation: At 0.05 significance level, we reject the null hypothesis ($p\text{-value}=0.009$) and conclude that the true birthweight means for smokers and non-smokers are significantly different.

One-Sample T-test

- Hypothesis to be tested (two-sided):

$$H_0: \mu = \mu_0 \text{ vs } H_1: \mu \neq \mu_0$$

Example: is the average age of mother's in this population different from 26?

R function:

```
t.test(mydata$myvar, mu = mu_null)
```

Options:

- Default alternative = ‘two-sided’
 - Can be changed to alternative = ‘less’ or alternative = ‘greater’
- Default alpha = 0.05

One-Sample T-test

Outcome:
Continuous

Example: is the average age of mother's in this population different from 26?

```
> t.test(lowbirth2$age, mu=26)
```

```
One Sample t-test
```

```
data: lowbirth2$age
t = -7.1659, df = 188, p-value = 1.707e-11
alternative hypothesis: true mean is not equal to 26
95 percent confidence interval:
 22.47779 23.99840
sample estimates:
mean of x
 23.2381
```

95% CI: (22.48, 23.99).

We are 95% confident that the true mean mother's age is b/w approximately 23 and 24 yrs.

Interpretation: At 0.05 significance level, we reject the null hypothesis (p-value <0.0001) and conclude that the true mean mother's age is not equal to 26.

Analysis of Variance (ANOVA)

- Use to compare the (continuous) outcomes across 3 or more groups
- Model assumptions:
 - Independent samples
 - Responses within the groups are independent and identically distributed (i.i.d)
 - Residuals are normally distributed
 - Equality of variances across groups

Analysis of Variance (ANOVA)

Example: is the birthweight of babies significantly different by race?

R function:

```
lm(cont_outcome~cat_predictor, data = mydata)
```

Notes:

- R also has an `aov()` function, but `lm()` is broader including linear regression models
- Better to declare the categorical variable/predictor as a *factor*, o/w it will be considered a continuous measurement.

Analysis of Variance (ANOVA)

Example: is the birthweight of babies significantly different by race?

```
> example1 = lm(bwt~race, data=lowbirth2)
> anova(example1)
Analysis of Variance Table

Response: bwt
            Df  Sum Sq Mean Sq F value    Pr(>F)
race          2 5070608 2535304  4.9719 0.007879 **
Residuals 186 94846445  509927
---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 '.' 1
>
```

Outcome: Continuous

Predictor: Categorical
(3 or more levels)

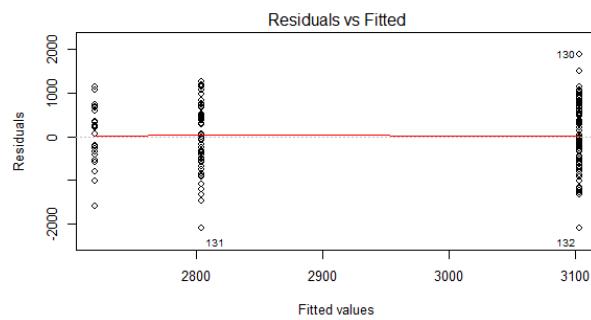
Interpretation: At 0.05 significance level, we reject the null hypothesis ($p\text{-value}=0.008$) and conclude that there is a significant difference in mean birthweight by race.

Next question: Where are these differences coming from? Try pairwise comparisons.
Need to adjust for multiple comparisons (Tukey, Bonferroni, Scheffe, etc.).
Topic covered in future courses!

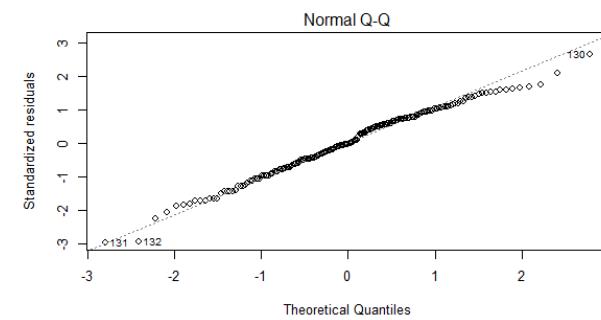
Analysis of Variance (ANOVA)

Example: checking model assumptions – plot(example1)

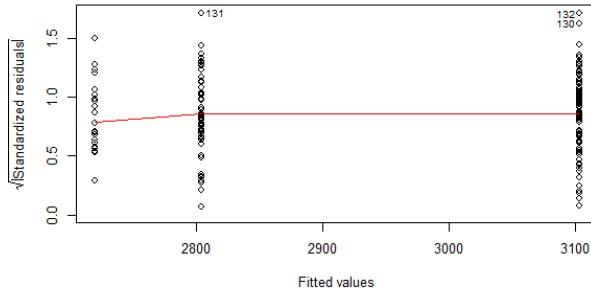
Constant Variance



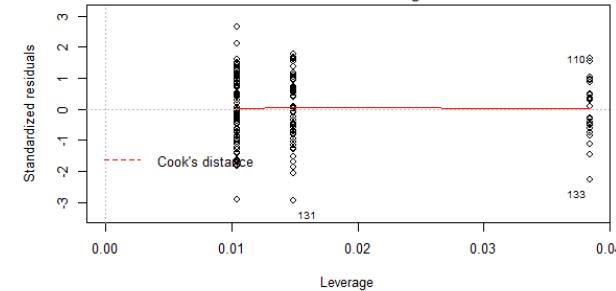
Normality of Residuals



Scale-Location



Residuals vs Leverage



Outliers/Influential Points

Categorical Data Analysis

- Categorical outcome (Y) with 2 levels (binary) or ≥ 3 levels (nominal or ordinal)
- Examples:
 - Nominal: race/ethnicity
 - Ordinal: clothing sizes (S, M, L, XL)
 - Binary: Disease/ No Disease; Republican/Democrat
- Predictor variables (X) can take on any form: binary, categorical, and/or continuous

Chi-Squared Test of Independence

- Use two categorical variables (row and column) to test whether they are independent or associated
- Hypotheses:

H_0 : variables A and B are independent

vs

H_1 : variables A and B are not independent

Test statistics: $\chi^2 = \sum \frac{(Observed - Expected)^2}{Expected}$

Chi-Squared Test of Independence

- Create a ($r \times c$) table
 - r represents the number of levels for the row variable
 - c represents the number of levels for the column variable
 - Most common example is a 2×2 table
- Use the observed and expected counts in each cell to calculate the chi-squared statistics
- If low expected cell counts (< 5), use Fisher's Exact test instead

Chi-Squared Test

- First you need to tabulate the two categorical variables and then apply chisq.test() to this table

R function:

```
table(mydata$row_var, mydata$col_var)  
chisq.test(mydata$row_var, mydata$col_var)
```

Example: is there an association between history of uterine irritability and having a low birthweight baby?

```
> library(MASS)  
> tb1 <- table(low_birth_all$ui, low_birth_all$low)  
> chisq.test(tb1)
```

Chi-Squared Test

Outcome: Categorical
Predictor: Categorical

Example: is there an association between smoking and having a low birthweight baby?

```
> table(lowbirth$smoke, lowbirth2$low)

  0  1
0 86 29
1 44 30
> chisq.test(lowbirth2$smoke, lowbirth2$low)

Pearson's Chi-squared test with Yates' continuity correction

data: lowbirth2$smoke and lowbirth2$low
X-squared = 4.2359, df = 1, p-value = 0.03958
```

Interpretation: At 0.05 significance level, we reject the null hypothesis ($p\text{-value}=0.040$) and conclude that there is a significant association between smoking and having a low birthweight baby.

Fisher's Exact Test

- Use instead of chi-squared test when low expected cell counts (<5)

Example: is there an association between smoking and having a low birthweight baby?

```
> fisher.test(lowbirth2$smoke, lowbirth2$low)

  Fisher's Exact Test for Count Data

data: lowbirth2$smoke and lowbirth2$low
p-value = 0.03618
alternative hypothesis: true odds ratio is not equal to 1
95 percent confidence interval:
 1.028780 3.964904
sample estimates:
odds ratio
 2.014137
```

Interpretation: At 0.05 significance level, we reject the null hypothesis ($p\text{-value}=0.036$) and conclude that there is a significant association between smoking and having a low birthweight baby.

Thank you!

Contact me: cmm2212@cumc.columbia.edu

Useful Resources:

[http://p8105.com/topic visualization and eda.html](http://p8105.com/topic_visualization_and_eda.html)

<https://stats.idre.ucla.edu/r/>

Visit our BERD EDU website for additional resources:

http://irvinginstitute.columbia.edu/resources/biostat_educational_initiatives.html